



MEINF

HIGH PERFORMANCE COMPUTING

---

# HPC Project: Heat Diffusion Equation

Second Delivery (Hybrid Implementation with OpenMP-MPI)

---

MARC GASPÀ JOVAL  
RAUL HIDALGO CABALLERO

May 9, 2025

# Contents

<b>1</b>	<b>Strategy</b>	<b>3</b>
1.1	Static mapping of tasks . . . . .	3
1.2	Communication . . . . .	4
1.3	Load Balancing . . . . .	5
<b>2</b>	<b>Scalability of the Program</b>	<b>6</b>
2.1	Performance Measurement Methodology . . . . .	6
2.2	Performance Metrics and Their Formulas . . . . .	6
2.3	Speedup and Efficiency . . . . .	8
2.4	Overhead Analysis . . . . .	10

## List of Tables

1	Execution times of the <code>heat_serial</code> program. . . . .	6
2	Execution times of the <code>heat_parallel</code> hybrid program with 2 nodes. . . .	7
3	Execution times of the <code>heat_parallel</code> hybrid program with 4 nodes. . . .	7
4	Execution times of the <code>heat_parallel</code> hybrid program with 5 nodes. . . .	7
5	Speedup of the parallel program with respect to the serial one for different thread counts. . . . .	8
6	Efficiency of the parallel program with respect to the serial one for different thread counts. . . . .	9
7	Overhead percentage of the parallel program for different thread counts and matrix sizes (normalized). . . . .	11

# 1 Strategy

From the two possible strategies, we used the **Static mapping of tasks** strategy. The main reason we chose this approach is that it is faster to implement and easier to understand.

## 1.1 Static mapping of tasks

In our case, we have a 2D grid of size  $N \times N$ , and we want to divide it into  $P$  subgrids, which correspond to the number of slots we will allocate in the cluster.

Although the image in the assignment suggests dividing into square subgrids, we decided to divide the grid into rectangular subgrids. This choice allows for better memory locality and improved cache usage.

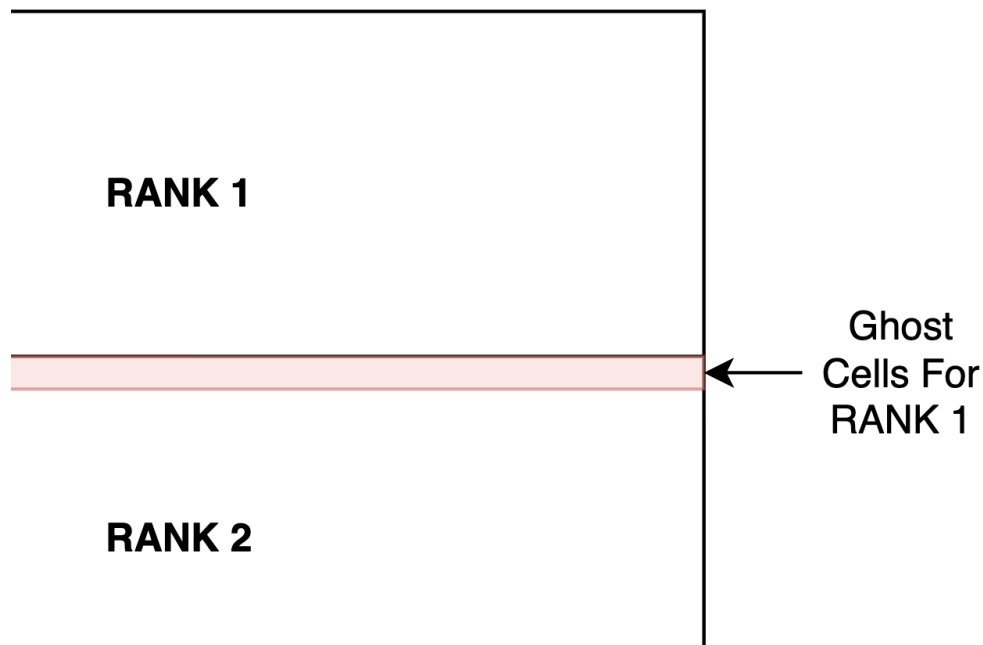


Figure 1: Illustration of the grid division into rectangular subgrids. And the ghost cells that are used to communicate with the neighbors.

## 1.2 Communication

In our implementation, communication is performed at three key stages:

- **Initialization:** Each process must receive its portion of the global grid, including two extra rows (ghost rows) for neighbor data. This is achieved using `MPI_Scatterv`, which distributes variable-sized blocks of data to each process.
- **Iteration:** At the beginning of each iteration, processes exchange boundary rows with their neighbors to update ghost rows. This is done using `MPI_Isend` and `MPI_Irecv` for non-blocking communication, followed by `MPI_Waitall` to ensure all communications complete before proceeding with computation.
- **Finalization:** Once computation is complete, each process sends its subgrid back to the master process. We use `MPI_Gatherv` to collect variable-sized subgrids and reconstruct the full grid in the master process before writing it to a file.

### 1.3 Load Balancing

Load balancing is not a concern in our case since the number of iterations and the number of grid cells are fixed. To ensure an even workload distribution, we divide the grid into rectangular subgrids of nearly equal size among the processes.

One limitation is that we do not currently handle the case where the grid size is not divisible by the number of processes. In such scenarios, the last part of the grid may lose some rows. This could be improved by implementing a remainder handling strategy in the grid division.

## 2 Scalability of the Program

### 2.1 Performance Measurement Methodology

In this case, as we are using MPI, we will use the `MPI_Wtime` function to measure the time taken by each process.

### 2.2 Performance Metrics and Their Formulas

Matrix Size	Steps			
	100	1000	10000	100000
$100 \times 100$	0.010000s	0.100000s	1.080000s	10.820000s
$1000 \times 1000$	1.200000s	12.780000s	119.870000s	1070.680000s
$2000 \times 2000$	4.750000s	47.630000s	503.060000s	4310.300000s

Table 1: Execution times of the `heat_serial` program.

---

Nodes: 2	Steps			
	100	1000	10000	100000
$1000 \times 1000$	3.226295s	28.773441s	269.608456s	2451.152783s
$100 \times 100$	3.008608s	30.182583s	151.455907s	997.242430s
$2000 \times 2000$	4.463269s	42.416510s	323.633867s	2576.611911s

---

Table 2: Execution times of the `heat_parallel` hybrid program with 2 nodes.

---

Nodes: 4	Steps			
	100	1000	10000	100000
$1000 \times 1000$	3.466167s	32.299433s	318.805040s	3184.036661s
$100 \times 100$	2.661267s	26.713319s	372.033094s	2572.231210s
$2000 \times 2000$	5.508030s	47.554560s	437.503118s	3524.018770s

---

Table 3: Execution times of the `heat_parallel` hybrid program with 4 nodes.

---

Nodes: 5	Steps			
	100	1000	10000	100000
$1000 \times 1000$	2.874023s	27.269382s	382.538333s	3183.711562s
$100 \times 100$	2.222764s	23.600960s	310.150252s	3098.737703s
$2000 \times 2000$	4.407035s	36.218476s	428.217962s	3336.180224s

---

Table 4: Execution times of the `heat_parallel` hybrid program with 5 nodes.



## 2.3 Speedup and Efficiency

Speedup formula:

$$S = \frac{T_{serial}}{T_{parallel}} \quad (1)$$

Matrix Size	Threads	Steps			
		100	1000	10000	100000
100 × 100	2	0.003324	0.003313	0.007131	0.010850
	4	0.003758	0.003743	0.002903	0.004206
	5	0.004499	0.004237	0.003482	0.003492
1000 × 1000	2	0.371944	0.444160	0.444608	0.436807
	4	0.346204	0.395673	0.375998	0.336265
	5	0.417533	0.468657	0.313354	0.336299
2000 × 2000	2	1.064242	1.122912	1.554411	1.672856
	4	0.862377	1.001586	1.149843	1.223121
	5	1.077822	1.315075	1.174776	1.291987

Table 5: Speedup of the parallel program with respect to the serial one for different thread counts.

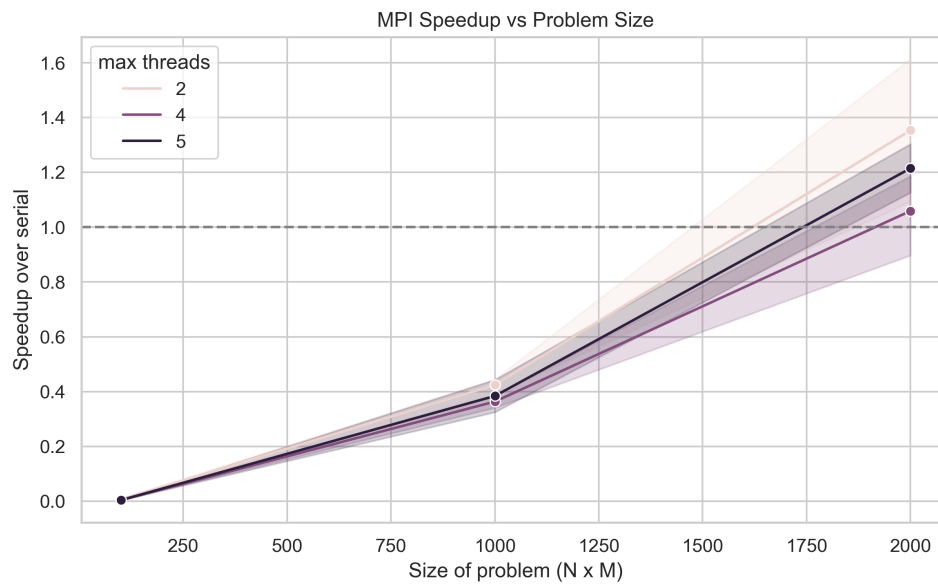


Figure 2: Scalability analysis of the program for different matrix sizes and thread counts.

Efficiency formula:

$$E = \frac{S}{P} \quad (2)$$

Here,  $P$  represents the number of processes utilized. We used 2, 4, and 5 processes for our tests.

Matrix Size	Threads	Steps			
		100	1000	10000	100000
$100 \times 100$	2	0.001662	0.001657	0.003565	0.005425
	4	0.000939	0.000936	0.000726	0.001052
	5	0.000900	0.000847	0.000696	0.000698
$1000 \times 1000$	2	0.185972	0.222080	0.222304	0.218403
	4	0.086551	0.098918	0.093999	0.084066
	5	0.083507	0.093731	0.062671	0.067260
$2000 \times 2000$	2	0.532121	0.561456	0.777205	0.836428
	4	0.215594	0.250397	0.287461	0.305780
	5	0.215564	0.263015	0.234955	0.258397

Table 6: Efficiency of the parallel program with respect to the serial one for different thread counts.

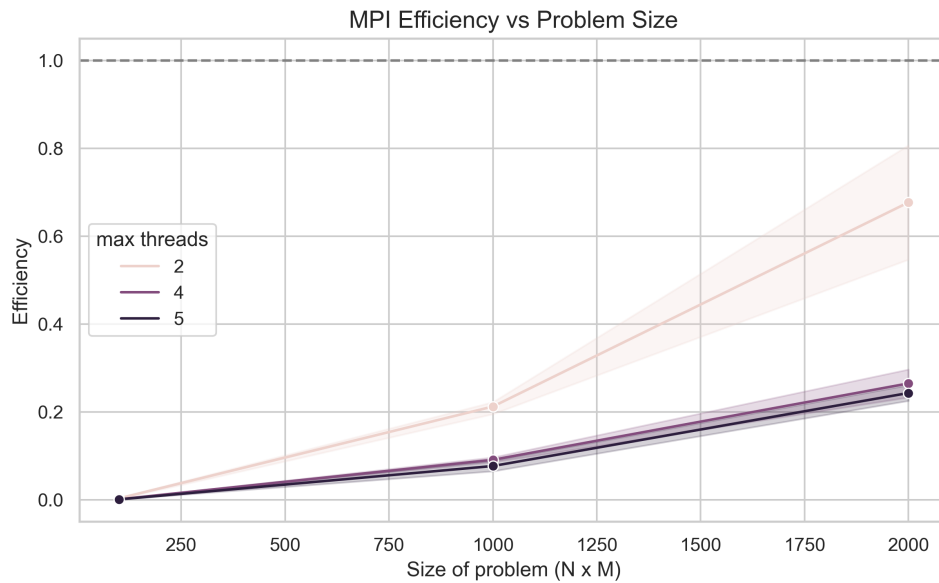


Figure 3: Efficiency analysis of the program for different matrix sizes and thread counts.

## 2.4 Overhead Analysis

As you can see in the previous data, the overhead in our implementation is important, only seeing speedups over 1 after the 2000x2000 matrix. This is because the overhead of the communication is greater than the time spent in the computation.

We will calculate the overhead using the following formula:

$$O = \frac{T_{parallel} \cdot P - T_{serial}}{T_{parallel} \cdot P} \cdot 100 \quad (3)$$

Which is normalized to 100% and gives us the percentage of time spent in communication.

Matrix Size	Threads	Steps			
		100	1000	10000	100000
$100 \times 100$	2	0.060072	0.602652	3.018318	19.836649
	4	0.106351	1.067533	14.870524	102.781048
	5	0.111038	1.179048	15.496713	154.828685
$1000 \times 1000$	2	0.005253	0.044767	0.419347	3.831626
	4	0.012665	0.116418	1.155350	11.665467
	5	0.013170	0.123567	1.792822	14.847878
$2000 \times 2000$	2	0.002088	0.018602	0.072104	0.421462
	4	0.008641	0.071294	0.623476	4.892888
	5	0.008643	0.066731	0.819015	6.185301

Table 7: Overhead percentage of the parallel program for different thread counts and matrix sizes (normalized).

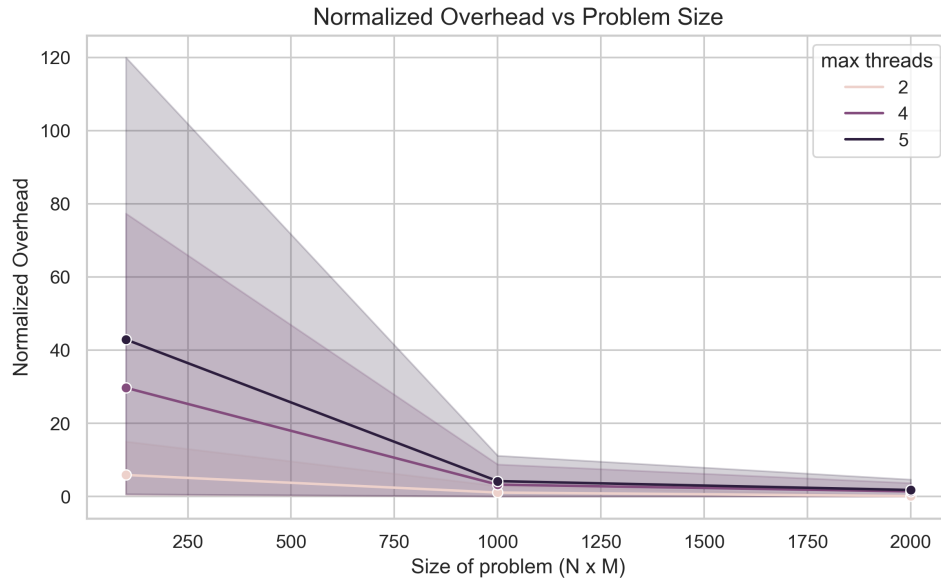


Figure 4: Overhead analysis of the program for different matrix sizes and thread counts.

As we can see in the table and image, the overhead is very high for small matrix sizes, but it decreases as the matrix size increases. This is because the time spent in communication is less than the time spent in computation.