



MEINF

HIGH PERFORMANCE COMPUTING

---

# HPC Project: Heat Diffusion Equation

First Delivery (OpenMP)

---

MARC GASPÀ JOVAL  
RAUL HIDALGO CABALLERO

April 11, 2025

# Contents

<b>1</b>	<b>Main Decisions Taken to Parallelize the Problem</b>	<b>3</b>
1.1	Grid Initialization (initialize_grid) . . . . .	3
1.1.1	Parallelizable Section . . . . .	3
1.1.2	Strategy . . . . .	3
1.2	Computation of the Heat Equation (solve_heat_equation) . . . . .	4
1.2.1	Parallelizable Section . . . . .	4
1.2.2	Strategy . . . . .	4
1.3	Data Dependency Management . . . . .	5
1.3.1	Overall Approach . . . . .	5
1.3.2	Synchronization . . . . .	5
1.4	Serial Output Operation . . . . .	5
1.4.1	Analysis . . . . .	5
1.4.2	Decision . . . . .	5
<b>2</b>	<b>Scalability of the Program</b>	<b>6</b>
2.1	Performance Measurement Methodology . . . . .	6
2.2	Performance Metrics and Their Formulas . . . . .	6
2.3	Speedup and Efficiency . . . . .	7
2.4	Scalability Analysis . . . . .	7

## List of Tables

1	Execution times of the <code>heat_serial</code> program. . . . .	6
2	Execution times of the <code>heat_parallel</code> program. . . . .	6
3	Speedup of the parallel program with respect to the serial one. . . . .	7
4	Efficiency of the parallel program with respect to the serial one. . . . .	7

# 1 Main Decisions Taken to Parallelize the Problem

We encountered we SMPD (Single Program Multiple Data) problem which we could parallelize using OpenMP.

This section details each part of the serial code that was identified as suitable for parallelization, along with the strategies employed.

Also we applied some optimizations to the code, like not recalculating the array indexes for each operation, instead we calculated them once and stored them in a variable.

## 1.1 Grid Initialization (`initialize_grid`)

### 1.1.1 Parallelizable Section

The initialization loop sets the temperature for each cell independently.

```
for (i = 0; i < nx; i++) {  
    for (j = 0; j < ny; j++) {  
        int inyj = i * ny + j;  
        if (i == j) { grid[inyj] = 1500.0; }  
        else if (i == nx - 1 - j) { grid[inyj] = 1500.0; }  
        else { grid[inyj] = 0.0; }  
    }  
}
```

### 1.1.2 Strategy

In this case as there aren't any dependencies we used `#pragma omp parallel for` with the `collapse(2)` clause on the nested loops to distribute the iterations.

## 1.2 Computation of the Heat Equation (solve\_heat\_equation)

### 1.2.1 Parallelizable Section

In this case, as one step is dependent of the previous one we cannot parallelize the outer loop. However, the inner loop can parallelize two parts, the first one is the application of the heat equation:

```
for (i = 1; i < nx - 1; i++) {
    for (j = 1; j < ny - 1; j++) {
        int inyj = i * ny + j;
        new_grid[inyj] = grid[inyj]
            + r * (grid[(i + 1) * ny + j] + grid[(i - 1) * ny + j] - 2 * grid[inyj])
            + r * (grid[inyj + 1] + grid[inyj - 1] - 2 * grid[inyj]);
    }
}
```

The second one is the application of the boundary conditions:

```
for (i = 0; i < nx; i++) {
    new_grid[0 * ny + i] = 0.0;
    new_grid[ny * (nx - 1) + i] = 0.0;
}

for (j = 0; j < ny; j++) {
    new_grid[0 + j * nx] = 0.0;
    new_grid[(ny - 1) + j * nx] = 0.0;
}
```

### 1.2.2 Strategy

In this case we really used a similar solution for both cases, as each cell can be calculated independently from the others.

We used `#pragma omp parallel for` with the `collapse(2)` clause on the heat equation loop to distribute the iterations. We defined the variables `i` and `j` as private.

For the boundary conditions we added `#pragma omp parallel for` to the loops that set the values of the boundaries. We also defined the variable `i` as private in the first loop and `j` in the second one.

## 1.3 Data Dependency Management

### 1.3.1 Overall Approach

A two-buffer scheme is employed throughout the computational phase. While one buffer holds the original data (read-only), the other is used for storing computed updates.

### 1.3.2 Synchronization

The swapping of pointers between `grid` and `new_grid` is performed outside parallel regions, ensuring data integrity for the subsequent computation cycle.

## 1.4 Serial Output Operation

### 1.4.1 Analysis

The generation and writing of the BMP file require a strict order of operations (bottom-to-top row order and specific padding for each row), which is highly sequential.

### 1.4.2 Decision

The output phase remains serial due to the order dependency and the relatively small impact on overall performance. But what we did is optimize the code without parallelizing it. Exactly we reduced the number of times we write to the file by writing three bytes at once instead.

## 2 Scalability of the Program

### 2.1 Performance Measurement Methodology

Execution times are measured using `omp_get_wtime()`. Testing is conducted using various grid sizes (matrix sizes) and numbers of iteration steps.

Just to clarify, we run the following tests in the moore cluster where we used the full 4 cores of each node.

### 2.2 Performance Metrics and Their Formulas

Matrix Size	Steps			
	100	1000	10000	100000
$100 \times 100$	0.010000s	0.100000s	1.080000s	10.820000s
$1000 \times 1000$	1.200000s	12.780000s	119.870000s	1070.680000s
$2000 \times 2000$	4.750000s	47.630000s	503.060000s	4310.300000s

Table 1: Execution times of the `heat_serial` program.

Matrix Size	Steps			
	100	1000	10000	100000
$100 \times 100$	0.006104s	0.040334s	26.465748s	33.893508s
$1000 \times 1000$	1.671250s	13.486841s	33.677816s	698.878759s
$2000 \times 2000$	5.421353s	33.636055s	135.988130s	1499.213647s

Table 2: Execution times of the `heat_parallel` program.

## 2.3 Speedup and Efficiency

Speedup formula:

$$S = \frac{T_{serial}}{T_{parallel}} \quad (1)$$

Matrix Size	Steps			
	100	1000	10000	100000
100 × 100	1.63	2.48	0.04	0.32
1000 × 1000	0.72	0.95	3.56	1.53
2000 × 2000	0.87	1.42	3.69	2.88

Table 3: Speedup of the parallel program with respect to the serial one.

Efficiency formula:

$$E = \frac{S}{P} \quad (2)$$

Where  $P$  is the number of threads used. In our case, we used 4 threads.

Matrix Size	Steps			
	100	1000	10000	100000
100 × 100	0.41	0.62	0.01	0.08
1000 × 1000	0.18	0.24	0.89	0.38
2000 × 2000	0.22	0.36	0.92	0.72

Table 4: Efficiency of the parallel program with respect to the serial one.

## 2.4 Scalability Analysis

The program demonstrates good scalability, particularly for larger matrix sizes and higher iteration counts. The speedup and efficiency metrics indicate that the parallel implementation effectively utilizes available resources, especially for larger problem sizes.